

Juan Manuel Fraga Sastrías · ORCID orcid.org/0000-0002-9255-278X · www.docfraga.com/en/whitepapers/gx10-17-dias · License CC BY 4.0

AI infrastructure July 22, 2026 · 17 min read

From Amazon to production in 17 days

I bought a 128 GB AI workstation and in 17 days it went from the box to serving a 35-billion-parameter model in production for my agents. Here are the real numbers, the exact inference-engine setup, and the technical gotchas nobody documents —enough for you to decide, and for your technical advisor to confirm it's done right.

My conversational assistants —the ones answering over WhatsApp— ran inference on a Mac. It worked, until it didn't: as I added agents, a single machine choked and latency spiked at peak hours. I needed **dedicated, on-premises AI compute**, no cloud. Not for the hype: for **data privacy** (health information that shouldn't leave my network), recurring cost (stop paying per token), and control (pick the model and fine-tune it myself). This is the honest log of how I solved it, the good parts and the uncomfortable ones.

The machine

An **ASUS Ascent GX10**, a DGX Spark: an **NVIDIA GB10 Grace Blackwell** superchip with **128 GB of unified memory** (about 121 GB usable by models), 20 ARM cores, and an Ubuntu-family OS. Price: around **\$88,000 MXN** —what a premium laptop costs.

The detail a technical reader will want: the GB10 is so new it reports **compute capability sm_121**, which forces the **CUDA 13** branch (with PyTorch built against it). Its known limit isn't raw compute but **memory bandwidth (~270 GB/s)**: that's what caps text generation speed. So it shines at **sustained multi-user inference**—many conversations at once—and doesn't pretend to beat a data-center server on extreme bursts. For serving agents locally, that's exactly the curve I wanted.

The 17 days

- **Day 0.** Bought it online. The same model went up ~7% in price days later: buying on time is engineering too.
- **Days 1–15.** In transit. The box crossed half the country.
- **Day 16.** Arrives. I tested it outside the rack first: clean boot, private network configured, key-based remote access ready. Idle: GPU at ~37 °C and 5 W.
- **Day 17.** Overnight —automatically, while I slept— it came up serving its first model, with the container set to auto-restart on any crash. By the next day it was answering a production agent.

The model

Qwen3.6-35B: a **mixture-of-experts** (MoE) model with 35 billion total parameters but only **~3 billion active per token** —it routes each token to 8 of its 256 experts plus one shared. That's why it's fast despite its size: it computes like a small model and "knows" like a big one. It's multimodal (text, image, and video), open-weight (Apache 2.0), and handles very long contexts (262K tokens, up to ~1M with extension).

One nuance that turned out to be decisive: it isn't a classic attention-only MoE but a **hybrid architecture** —it mixes attention layers with state-space layers of the **Mamba / Gated DeltaNet** family. That choice makes it efficient on long contexts, but changes the rules for how you serve it (I pay for that in gotcha 2). I run it with **vLLM** in a container, with weights quantized to 4 bits in the **NVFP4** format built for this GPU generation.

How I serve it: the tuned engine

Serving a model isn't "switching it on": it's a handful of decisions that separate a demo from something that holds real users. Here are the ones I turned on in vLLM and why, no hand-waving:

Decision	Why
NVFP4 (4-bit) quantization	This Blackwell generation's 4-bit format; the 35B model's process reserves ~34 GB (at 0.28 GPU memory utilization, KV cache included) with the MARLIN NvFp4 MoE backend.
FP8 KV cache	Doubles usable context memory: ~70 GiB of cache at 40K tokens → dozens of live conversations at once.
Prefix caching + chunked prefill	Reuses the shared system prompt across turns and slices prefill so one long request doesn't block everyone else.
Pinned minimum batch (8192 tokens)	A hard requirement of the hybrid Mamba architecture: without it the engine won't even start (see gotcha 2).
Speculative decoding (MTP ×3)	The model drafts several tokens per step and verifies them in a block: lower per-response latency.
Tool-calling + reasoning parser	Lets agents invoke tools (the model's XML format) and separates the chain of thought from the final answer.

Together, these pieces light up the GPU's fast paths: the quantized MoE backend, the optimized attention engine, the compressed context cache, and speculative decoding. The measurable effect is twofold: it **lowers latency** for a single

response and, above all, **multiplies how many conversations fit in parallel** without degrading.

The numbers

Real run, model already warm and the setup optimized:

Concurrent conversations	Latency	Aggregate throughput
1	0.7 s	112 tok/s
4	1.1 s	271 tok/s
8	1.3 s	493 tok/s
16	1.8 s	669 tok/s

The interesting part isn't a single number, it's the **shape**: latency barely rises from 1 to 16 conversations (0.7 s → 1.8 s) while throughput multiplies ~6× (112 → 669 tokens/s). That's the signature of a GPU that **genuinely batches requests** — processes them together in one pass— rather than doing them in a queue. It's exactly what an Apple chip can't: when the second conversation arrives, it serializes. Here they coexist.

After tuning, a single-user response dropped from **2.9 s to 0.7 s**. During the test: 95% GPU utilization, 57 °C, and just 28 W —less than a light bulb. And by the cache math: with the KV in FP8, at 40K tokens of context ~70 GiB are left for live conversations, a theoretical ceiling of dozens of simultaneous users. In real operation I reserve less memory on purpose, to leave room for the other services (below).

The gotchas nobody documents

1 • A GPU so new it breaks the recipes

The GB10 (compute capability sm_121) requires CUDA 13 and the newest inference-server image. On an older version the NVFP4 weight loader blows up looking for parameters that don't exist in that format, and the model *won't even load*. No guide warned you yet: you have to demand the new image, not the "stable" one from a few months ago.

2 • It isn't a "normal" MoE: the Mamba toll

Because it's hybrid (Mamba-style state layers), turning on prefix caching makes the engine *align* memory in fixed-size blocks. If the maximum token batch falls below that block, the engine *won't even start* and throws a cryptic assertion error. The fix is to pin the minimum batch high (8192 tokens). A single parameter, in no tutorial, and—worse—it looked like a memory problem when it wasn't. It cost me hours chasing the wrong lead.

3 • A dirty power cut corrupts the container

After an unclean outage, the container runtime's build layer corrupts and the service won't come up. The fix is to move the damaged layer aside, clear the failed service state, and restart the runtime; containers marked to revive come back on their own. A real-world resilience lesson, not a lab one: on your own infrastructure, the power cut is part of the failure model, and you design so the machine returns by itself.

Three workloads in one box

The same machine doesn't serve one model, it serves three workloads at once, each for a different job:

General chat

CT segmentation

Qwen3.6-35B, resident. It's the one serving the conversational agents.

A 3D vision model that delineates ~127 anatomical structures in a CT study in seconds. It runs on demand, not resident.

Search embeddings

The vectorizer that turns documents into numbers for the agents' semantic search (RAG).

I also tried a fourth, a medical text model (gpt-oss-20b and MedGemma-27B), and discarded it: the general Qwen beat it in my internal evaluation, and MedGemma, being dense, topped out at ~13 tokens/s. Production runs Qwen3.6-35B; the machine is also used from time to time for tests with other models, and those tests are done by stopping the production model first (unified memory is shared; never in parallel): the video-generation test ran at night with Qwen stopped. The remaining three don't all fit at full size in 121 GB, so there's a **memory choreography**: by day I reserve just enough for the agents and leave room for CT segmentation to enter in bursts; by night I shut down chat and free the machine to fine-tune models and build the search indexes. The same box works a day shift (serving) and a night shift (improving), without buying a second one.

How I choose what runs: my own arenas

Public benchmarks don't answer my questions: they don't tell me whether a model can handle my cases, my tools or my language. So I build **my own arenas**: I pit models against each other on my own cases —clinical cases with a reference answer, tool chains from my agents with a deliberate trap, Mexican-Spanish turns— and blind judges score them on criteria I define: task success, tool-use safety, calibration, and full-turn latency. That is how it was decided what runs in production and what gets dropped, including the medical model above.

Two lessons. First: **tokens per second don't predict the experience**; a dense model with explicit reasoning was only 2.4× slower per token, but took 71 s per turn against 2.2 s. Second: a model can be good on average and still unsafe with irreversible actions; the production one executed a deletion it shouldn't have in 6 of 10 repetitions. Swapping models only moves the probability, so the answer was architectural: confirmation gates in code and an irreversibility classifier that escalates those turns to a cloud model. The arenas, with their figures, are in the technical appendix.

In production

The model doesn't connect directly to each agent: it goes through an **in-house gateway** that speaks the industry standard (the same API contract the cloud uses), so for the agent, talking to my local GPU or to a remote model is identical—only a routing string changes. The gateway **routes per agent** (each assistant can have its own model) and, if the local one fails, it *falls back* automatically to a cloud model with a **circuit breaker** that waits a few minutes before retrying the local one, so it doesn't hammer a downed service. The agents never see the error; they see a response. That's the difference between an experiment and something you rely on for work.

Technical appendix

For anyone who needs to reproduce or audit this work: the parameters and measurements behind the article. Only what was recorded is listed; anything not measured is stated as such.

1 • Memory budget on 128 GB of unified memory

That is 128 GB physical, ~121 GB usable. CPU and GPU share one pool, so a "CPU RAM offload" frees nothing: a video model's official repository weighed 123.6 GB and didn't fit even running alone, so quantizing was a requirement, not an option.

Component	Reservation or measurement	Note
Qwen3.6-35B NVFP4 (vLLM, production)	~34 GB reserved (<code>--gpu-memory-utilization 0.28</code>)	Fixed 7.0 GiB FP8 KV cache; 128K context.
bge-m3 embeddings	~1.6 GB	1024 dimensions; resident.
Dense 31B test model (stopped on 2026-09-12)	~30 GB reserved (25 %) · 4.7 GB actually used	10 requests in 10 days, all local.
Machine state, 2026-09-12	93 of 121 GB in use (2 free) → 48 in use (73 available) after stopping the test model	Freed ~45 GB, more than the reservation: the process carries overhead.
Batch video generation (MiniMax-H3, pruned ~42.5 GB weights)	Total peak 92 GB with Qwen stopped · 115 of 121 GB with Qwen running	15 s clip: 989 s vs 1300 s; 5 s clip: 241 s vs 370 s (~35 % slower from memory pressure).

- **Reserving is not using.** `--gpu-memory-utilization` sets memory aside; it doesn't reflect consumption. Plan with what is reserved. Lowering the reservation is sometimes enough to gain headroom without stopping anything.
- **A second vLLM competes for the same memory.** Starting it next to the production one killed the first one's engine and left it down ~6 min. Rule: stop first, never in parallel, and verify with a real inference, because the port kept answering with the engine dead.
- **Heavy jobs run at night with chat stopped.** A script restores chat and verifies it with a real inference before finishing.

2 • Benchmark methodology

Aspect	Record
Date and state	Jul 15, 2026, hot run (model already loaded and serving). Cold start was not measured.
Path	Requests through a private-network tunnel, from a separate machine, to the OpenAI-compatible endpoint of the vLLM container.

Aspect	Record
Configuration	<code>--gpu-memory-utilization 0.8</code> , <code>--max-model-len 40960</code> , <code>--max-num-seqs 32</code> , FP8 KV cache, MTP ×3.
Concurrency	1, 4, 8 and 16 parallel conversations.
What was measured	Per-response latency (s) and aggregate generation throughput (tokens/s) summed across the conversations.
Effect of tuning	Single user: 2.9 s → 0.7 s. During the run: GPU at 95 %, 57 °C, 28 W.
Not recorded	Load tool, prompt and output length, number of repetitions, and percentiles (p50/p95). The figures are point values, not distributions.

Second series, 2026-07-16, with a different configuration (`--gpu-memory-utilization 0.40`). Not comparable point by point with the “The numbers” table.

Parallel conversations	Aggregate throughput
1	92 tok/s
8	313 tok/s
16	407 tok/s
32	546 tok/s

At 32 conversations latency was around 5 s (~6× single-user throughput). The 1.11 M-token cache gives ~27 simultaneous requests at a full 40K context and up to 32 with normal prompts.

3 • vLLM configuration

GPU container with `--gpus all --runtime nvidia --ipc=host --restart unless-stopped`. The table in the body explains the general why; this one lists the exact values.

Parameter	Value	Why
Image	<code>vllm/vllm-openai:latest</code> (≥ 0.25 ; ~0.25.1 at the time)	With <code>v0.18.0-cu130</code> the NVFP4 loader fails (<code>KeyError ... experts.w2_input_scale</code>). The GB10 reports compute capability 12.1 (sm_121) and requires CUDA 13.
Model	<code>nvidia/Qwen3.6-35B-A3B-NVFP4</code>	Active backends: MARLIN NvFp4 MoE, FlashInfer and DeepGEMM.
<code>--gpu-memory-utilization</code>	0.6 $\rightarrow$ 0.8 (benchmark) $\rightarrow$ 0.40 $\rightarrow$ 0.28 (production)	A reservation, not usage. Lowered to coexist with other services on the same memory.
<code>--max-model-len</code>	32768 $\rightarrow$ 40960 $\rightarrow$ 131072 (2026-08-03)	Raised to 128K for a fallback coding tool that failed reading large files. It costs no memory (PagedAttention allocates by actual use); it only lowers theoretical max concurrency: 9.57 $\times$ at 40K $\rightarrow$ 4.04 $\times$ at 128K.
<code>--max-num-seqs</code>	4 $\rightarrow$ 32	Ceiling on simultaneous sequences.
<code>--max-num-batched-tokens</code>	8192	Hybrid Mamba architecture: with prefix caching it must be $\geq$ the block size (2144) or the engine won't start.
<code>--kv-cache-dtype</code>	fp8	Doubles usable context per GB: ~70 GiB of cache at 40K tokens in the benchmark configuration.

Parameter	Value	Why
<code>--enable-chunked-prefill · --enable-prefix-caching</code>	enabled	Shared system prompt reused across turns; chunked prefill keeps long prompts from blocking others.
<code>--speculative-config</code>	<code>{"method":"mtp","num_speculative_tokens":3}</code>	Speculative decoding using the model's own MTP heads.
<code>--tool-call-parser qwen3_xml · --enable-auto-tool-choice</code>	enabled	Tool use measured at 4/5 vs 2/5 on the Apple MLX server, which emits flat JSON and resolves only partially.
Reasoning	<code>--reasoning-parser qwen3</code> at first → removed; custom chat template with reasoning off by default and opt-in (<code>enable_thinking</code>)	With reasoning on by default the model spent all of <code>max_tokens</code> thinking and returned empty content; the parser also misfiled direct answers into the reasoning field.
Embeddings	<code>--runner pooling</code> (bge-m3, 1024 dim.)	<code>--task embed</code> does not exist in vLLM 0.25.1.

Tried and discarded

Model	Measurement	Cause or requirement
MedGemma-27B (dense)	BF16 4 tok/s · dynamic FP8 8 · int4 W4A16 Marlin + n-gram ~13; Qwen's MoE runs at ~100 tok/s.	head_dim 256/512 is unsupported by FlashInfer on sm_121, so it falls to a slow Triton kernel (vLLM issues #40677, #38887, #20865). Quantization speeds up weight GEMMs, not attention; n-gram speculation doesn't help free-form generation.

Model	Measurement	Cause or requirement
gpt-oss-20b (MXFP4)	34–44 tok/s; lost to Qwen in the internal evaluation (3/5 and 1/5 against 5/5).	Starting it needed <code>--enforce-eager</code> (CUDA-graph compilation hung ~40 min) and pre-downloaded tiktoken vocabularies (harmony's Rust client ignores the certificate store; vLLM #22525).

4 • Failures seen in production and their mitigation

Failure	Cause	Mitigation
Qwen engine killed when a second vLLM was started (2026-08-14, ~6 min out of service)	Shared unified memory: <code>CUDA error: an illegal memory access</code> . The utilization value of the container already running does not protect it.	Stop first, never in parallel; verify with a real inference.
<code>/health</code> returns 200 with the engine dead	The endpoint doesn't exercise the engine.	Probe with a real generation. Same pattern seen earlier on the Apple MLX server: the watchdog generates 5 tokens (10 s timeout) every 5 min and restarts up to twice before escalating to manual intervention.
Startup aborted by a Mamba assertion when relaunching the container	<code>--max-num-batched-tokens 8192</code> was omitted. The symptom misled: it looked like fragmentation OOM.	Copy every factory flag when relaunching any container.
Memory fragmentation after many container swaps	Load/unload churn in unified memory.	Reboot the machine; containers with <code>unless-stopped</code> revive by themselves.
Unclean power cut: the container service won't start	Corrupted Docker BuildKit database.	Move the damaged database aside, clear the failed service state, restart the service. The private-network client runs as a boot service (no graphical session), so the machine comes back on its own.

Failure	Cause	Mitigation
"Local engine unreachable" alert during planned overnight runs	The alert was legitimate: the engine was stopped on purpose.	Any overnight run must announce itself or silence the monitor for its window.
Models visible but unreadable inside the container	The Hugging Face cache stores symlinks; mounting only a subdirectory breaks them (<code>FileNotFoundError</code> even though the listing shows the model).	Build a tree of hardlinks, or mount the whole cache.
A fallback masking a broken primary (Jul–Aug 2026)	The model id carried a provider prefix and the real key rejected it with 403. Of all messages in the "local" channels, only one was served by the local model; 11 fallback events; with the breaker open nothing was logged.	Routing-metrics endpoint: configured vs who actually answered (<code>ok</code> / <code>divergent</code> / <code>no_data</code>); the log now records the reason (<code>primary_error</code> or <code>breaker_open</code>). Validated on 2026-09-05: a cloud key stopped being valid (401) and the detector flagged it.

5 • Routing: one gateway in front of local and cloud models

Agents talk to a single OpenAI-compatible endpoint. Moving from the local GPU to a cloud model means changing a model string in the channel configuration: the agent keeps its SDK and its contract, which is why none of them had to be rewritten.

Piece	Detail
Contract	One OpenAI-compatible gateway. Canonical form: OpenAI-style <code>image_url</code> ; litellm translates to Anthropic's format when the target is cloud.
Identity	One key per agent, with an allowed-models list and a requests-per-minute cap. The list is enforced at completion time (403), not when listing models. The model id goes without a provider prefix.
Primary and fallback	Four fields per channel (primary provider and model; fallback provider and model), with retry and a 10-minute breaker. A local model's fallback must be cloud. As of 2026-09-24: 19 channels — 11 with a local primary and Haiku fallback, 6 with a Haiku primary and 2 with Sonnet, both with a local fallback.

Piece	Detail
Assignment criterion	Measurement, not preference: over 624 generations. On simulation-education questions the local model fabricated citations in 16 of 20 answers, so those channels use a cloud primary.
Multimodal content	Images pass as <code>image_url</code> ; PDFs are rasterized (up to 20 pages, 150 dpi) for models without native PDF; audio is transcribed at the gateway by default (per-key policy: <code>transcribe</code> or <code>native</code>); plain text is spliced in (cap of 200,000 characters).
Sampling and reasoning	Cold sampling profile applied automatically when the turn carries tools; opt-in reasoning is forwarded only to local engines.
Gateway cost	~10 % for a single user: 107 tok/s direct vs 97 tok/s through the gateway.
Real traffic	30 days to 2026-07-28: 5,015 requests to the local engine, 15 to the cloud and 11 to the Apple MLX server (~99.5 % local).
Separate control plane	The gateway runs on a machine other than the GPU: three vLLM restarts in a single night had no impact on the fleet.

6 • Arenas run

The head-to-head model comparisons behind this article’s decisions. Figures are those recorded in each one; judges and rubrics are noted in each row.

Date	What was compared	Criterion	Decision and figures
2026-03-12	In-house tool to compare 2–4 Ollama models side by side.	23 prompts in 6 categories; TTFT, tokens/s, output tokens, total time, % of model on GPU and load time.	Basis for the later arenas; no decision attached.
2026-03-26 to 04-05	Oncology (clinical QA): local and cloud models, with and without RAG, with	15 clinical cases with a reference answer; blind judge (Claude Opus); 4 criteria:	They did not go to human evaluation while the local configurations weren’t competitive; retrieval

Date	What was compared	Criterion	Decision and figures
	and without fine-tuning, in rounds of 5, 6 and 8 configurations.	diagnostic accuracy, guideline adherence, completeness and clinical usefulness.	is reviewed before being re-enabled (see the preregistered study linked above). Only the setup is described here, not results.
2026-03-28	Agent tasks on local models (MLX). Heavy: Nemotron-3-Nano-30B, Mistral-Small-3.2-24B, Qwen2.5-Coder-32B, gpt-oss-20b. Light: Qwen3-VL-8B, Phi-4 4B.	Score per round (11 and 6 rounds for heavy, 7 text + 3 vision for light) and tokens/s.	Nemotron-3-Nano-30B: 37/55 (60–85 tok/s) vs 36/55, 35/55 and 30/55; second round 26/30, and Qwen3.5-27B (25/30) was dropped for latency (59–258 s). Light: Qwen3-VL-8B 29/35 on text with perfect vision, vs Phi-4 26/35 with no vision.
2026-04-02	Gemma 4 26B MoE, Gemma 4 31B, MiniMax M2.7, Nemotron 30B, Gemma 4 E4B and gpt-oss-20b.	6 multi-agent tasks, 2 rounds (12 battles), automatic LLM judge.	Averages 4.50 / 4.25 / 3.67 / 3.33 / 2.83 / 2.50. The 31B (14 tok/s) didn't justify its slowness against the 26B MoE; gpt-oss-20b was dropped (code review 1/5).
2026-04-12	Seven configurations for agents: Gemma 4 31B dense local, MiniMax via NIM and direct, Nemotron 30B, GLM5, Gemma 4B and Gemma 26B local.	2 rounds of 6 prompts (delegation, conversation, DevOps, logs, code review, strategy); Claude Sonnet judge, 1–5 scale; latency.	Gemma 31B 4.83 (~51 s), MiniMax via NIM 4.17, MiniMax direct 4.00, Nemotron 3.83, GLM5 3.50, Gemma 4B 3.50 (~16 s), Gemma 26B local 3.17. Gemma 31B became the local premium, Gemma 4B the fast responder, and the local 26B was removed.

Date	What was compared	Criterion	Decision and figures
2026-08-11	Tool use: Qwen3.6-35B-A3B (MoE, production) vs Muse Glimmer 30B (dense).	5 scenarios with tool chains from my agents and a deliberate trap (a destructive or forbidden tool); blind Claude Opus judge; 4-dimension rubric: tools, avoided the trap, arguments and Spanish.	Quality 60 vs 82/100; latency per turn 2.2 s vs 71 s (~130 vs ~1,750 tokens per turn; 80 vs 33 tok/s). The production model executed a record deletion in 6 of 10 repetitions (the other, 0 of 10). Glimmer was switched off; sensitive or irreversible turns escalate to a cloud model and confirmation gates are added in code.
2026-09-18	Ternary-Bonsai-2-27B (ternary weights, 1.72 bits) vs Qwen3.8-27B NVFP4 and Qwen3.6-35B-A3B.	75 generations, blind judges: tools (/20), traps, clinical (/20) and tok/s.	Tools 15.93 / 16.13 / 13.13; traps 5/15, 3/15, 6/15; clinical 15.60 / 16.80 / 16.60; 32.3 / 11.1 / 49.8 tok/s. It stays off: aggressive quantization costs variance (0/3 vs 2/3 on the record case, with a perfect 20/20 on the third) and calibration (an invented figure in a reassuring tone). The publish-without-approval case failed 9 of 9.
2026-09-28	Can this turn trigger an irreversible action? A commercial small-model classifier, a local encoder (GLiFormer) and	120 Mexican-Spanish turns (54 irreversible, 66 harmless; 60 commands, 40 look-alike harmless pairs, 20 negation or conditional);	Accuracy 98.3 / 98.3 / 82.5 %; false negatives 1 / 1 / 16; Brier 0.023 / 0.017 / 0.170; p50 307 / 148 / 30 ms. Calibration decided, not accuracy. Second round: Qwen with <code>logprobs</code>

Date	What was compared	Criterion	Decision and figures
	Qwen3.6 in production.	accuracy, false negatives, Brier and p50 latency.	(max_tokens=1 , temperature 0): 98.3 %, Brier 0.021, p50 91 ms; at threshold 0.3 it covers 92 % with 99.1 % correct and 0 false negatives. In production since 2026-09-28 at threshold 0.8: it escalates to a cloud model.

Local AI compute is no longer exotic. For the price of a premium laptop I have a workstation serving a 35B model to my agents, segmenting CT scans, and powering search, **without paying per token and without data leaving the premises**. And the honest part: the hard bit wasn't the money or the model. It was three configuration decisions —an updated image, a batch parameter, and an outage-recovery plan— that came in no guide. What's next: this same machine is where the retrieval (RAG) layer of my preregistered oncology study will be built and run, "**Clinically Structured Retrieval: Can a Well-Designed RAG Layer Add Value Over a Base Language Model in Clinical Oncology?.**"

Building something similar —for your clinic, your center, or your practice— and want to compare notes? **Write to me →**